

## EDS ACTIVITY NO.1

**NAME: SAKSHI BARAHTA**

**DIV: CS5**

**ROLL NO.: CS5-61**

**PRN: 202401100110**

**SUBJECT: EDS**

```
test.py x
C: > Users > Atulya > OneDrive > Desktop > test.py > ...
1 import kagglehub
2
3 # Download latest version
4 path = kagglehub.dataset_download("fnbalves/paper-reviews-data-set")
5
6 print("Path to dataset files:", path)
7
8 import pandas as pd
9 import numpy as np
10
11 # --- Create a Sample Dataset ---
12 # This is a dummy dataset to demonstrate the solutions.
13 # In a real scenario, you would load your data from a file (e.g., CSV).
14 # Make sure the column names in the code match your actual dataset columns.
15
16 data = {
17     'review_id': range(1, 21),
18     'paper_id': [101, 102, 101, 103, 102, 101, 103, 104, 102, 104, 105, 101, 103, 105, 102, 104, 101, 105, 103, 104],
19     'reviewer_id': [201, 202, 203, 201, 203, 202, 201, 204, 204, 203, 205, 201, 202, 205, 203, 204, 202, 205, 201, 203],
20     'review_text': [
21         'Good paper, clear methodology.',
22         'Needs significant improvements in writing.'
23     ]
24 }
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

Sample Data Head:

|   | review_id | paper_id | reviewer_id | ... | recommendation | reviewer_expertise | review_length |
|---|-----------|----------|-------------|-----|----------------|--------------------|---------------|
| 0 | 1         | 101      | 201         | ... | Accept         | High               | 30            |
| 1 | 2         | 102      | 202         | ... | Major Revision | Medium             | 42            |
| 2 | 3         | 101      | 203         | ... | Accept         | High               | 36            |
| 3 | 4         | 103      | 201         | ... | Minor Revision | Medium             | 42            |
| 4 | 5         | 102      | 203         | ... | Accept         | Low                | 35            |

[5 rows x 9 columns]

Sample Data Info:  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20 entries, 0 to 19

```

C: > Users > Atulya > OneDrive > Desktop > test.py > ...
58     'High', 'High', 'Medium', 'Low', 'Medium',
59     'High', 'Medium', 'High', 'Medium', 'Low'
60     ]
61 }
62
63 df = pd.DataFrame(data)
64
65 # Add a column for review length (number of characters)
66 df['review_length'] = df['review_text'].str.len()
67
68 # Introduce some missing values for demonstration
69 df.loc[[2, 5, 11], 'rating'] = np.nan
70 df.loc[[7, 14], 'reviewer_expertise'] = np.nan
71
72
73 print("--- Paper Review Data Analysis ---")
74 print("\nSample Data Head:")
75 print(df.head())
76 print("\nSample Data Info:")
77 df.info()
78 print("-" * 40)
79

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

Data columns (total 9 columns):

| # | Column             | Non-Null Count | Dtype          |
|---|--------------------|----------------|----------------|
| 0 | review_id          | 20 non-null    | int64          |
| 1 | paper_id           | 20 non-null    | int64          |
| 2 | reviewer_id        | 20 non-null    | int64          |
| 3 | review_text        | 20 non-null    | object         |
| 4 | rating             | 17 non-null    | float64        |
| 5 | submission_date    | 20 non-null    | datetime64[ns] |
| 6 | recommendation     | 20 non-null    | object         |
| 7 | reviewer_expertise | 18 non-null    | object         |
| 8 | review_length      | 20 non-null    | int64          |

dtypes: datetime64[ns](1), float64(1), int64(4), object(3)

memory usage: 1.5+ KB

## PROBLEM STATEMENTS AND ANSWERS 1 AND 2:

```
81
82 # Problem 1: Display the first 5 rows of the dataset.
83 print("\nProblem 1: Display the first 5 rows of the dataset.")
84 print("Solution:")
85 print(df.head())
86 print("-" * 40)
87
88 # Problem 2: Get a summary of the numerical columns in the dataset.
89 print("\nProblem 2: Get a summary of the numerical columns in the dataset.")
90 print("Solution:")
91 print(df.describe())
92 print("-" * 40)
93
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

Problem 1: Display the first 5 rows of the dataset.

Solution:

|   | review_id | paper_id | reviewer_id | ... | recommendation | reviewer_expertise | review_length |
|---|-----------|----------|-------------|-----|----------------|--------------------|---------------|
| 0 | 1         | 101      | 201         | ... | Accept         | High               | 30            |
| 1 | 2         | 102      | 202         | ... | Major Revision | Medium             | 42            |
| 2 | 3         | 101      | 203         | ... | Accept         | High               | 36            |
| 3 | 4         | 103      | 201         | ... | Minor Revision | Medium             | 42            |
| 4 | 5         | 102      | 203         | ... | Accept         | Low                | 35            |

[5 rows x 9 columns]

Problem 2: Get a summary of the numerical columns in the dataset.

Solution:

|       | review_id | paper_id  | reviewer_id | rating    | submission_date     | review_length |
|-------|-----------|-----------|-------------|-----------|---------------------|---------------|
| count | 20.00000  | 20.00000  | 20.000000   | 17.000000 | 20                  | 20.000000     |
| mean  | 10.50000  | 102.80000 | 202.750000  | 3.470588  | 2023-03-06 03:36:00 | 34.950000     |
| min   | 1.00000   | 101.00000 | 201.000000  | 2.000000  | 2023-01-15 00:00:00 | 21.000000     |
| 25%   | 5.75000   | 101.75000 | 201.750000  | 3.000000  | 2023-02-08 18:00:00 | 30.000000     |
| 50%   | 10.50000  | 103.00000 | 203.000000  | 4.000000  | 2023-03-12 12:00:00 | 35.500000     |
| 75%   | 15.25000  | 104.00000 | 204.000000  | 4.000000  | 2023-04-02 00:00:00 | 42.000000     |
| max   | 20.00000  | 105.00000 | 205.000000  | 5.000000  | 2023-04-15 00:00:00 | 45.000000     |
| std   | 5.91608   | 1.43637   | 1.409554    | 1.067570  | NaN                 | 7.570545      |

## PROBLEM STATEMENTS AND ANSWERS 3, 4 AND 5:

```
C:\Users\Atulya\OneDrive\Desktop> test.py / ...
94  # Problem 3: Find the total number of reviews in the dataset.
95  print("\nProblem 3: Find the total number of reviews in the dataset.")
96  print("Solution:")
97  total_reviews = len(df)
98  print(f"Total number of reviews: {total_reviews}")
99  print("-" * 40)
100
101  # Problem 4: Count the number of unique papers reviewed.
102  print("\nProblem 4: Count the number of unique papers reviewed.")
103  print("Solution:")
104  unique_papers = df['paper_id'].nunique()
105  print(f"Number of unique papers reviewed: {unique_papers}")
106  print("-" * 40)
107
108  # Problem 5: Count the number of unique reviewers.
109  print("\nProblem 5: Count the number of unique reviewers.")
110  print("Solution:")
111  unique_reviewers = df['reviewer_id'].nunique()
112  print(f"Number of unique reviewers: {unique_reviewers}")
113  print("-" * 40)
114
```

PROBLEMS    OUTPUT    DEBUG CONSOLE    TERMINAL    PORTS

Problem 3: Find the total number of reviews in the dataset.

Solution:

Total number of reviews: 20

-----

Problem 4: Count the number of unique papers reviewed.

Solution:

Number of unique papers reviewed: 5

-----

Problem 5: Count the number of unique reviewers.

Solution:

Number of unique reviewers: 5



## PROBLEM STATEMENTS AND ANSWERS 6 AND 7:

C: > Users > Atulya > OneDrive > Desktop > test.py > ...

```
115 # Problem 6: Calculate the average rating given across all reviews (excluding NaNs).
116 print("\nProblem 6: Calculate the average rating given across all reviews (excluding NaNs).")
117 print("Solution:")
118 average_rating = df['rating'].mean()
119 print(f"Average rating: {average_rating:.2f}")
120 print("-" * 40)
121
122 # Problem 7: Find the paper with the highest average rating.
123 print("\nProblem 7: Find the paper with the highest average rating.")
124 print("Solution:")
125 # Calculate average rating per paper, drop papers with no valid ratings
126 avg_rating_per_paper = df.groupby('paper_id')['rating'].mean().dropna()
127 # Find the paper(s) with the maximum average rating
128 if not avg_rating_per_paper.empty:
129     highestRatedPaperId = avg_rating_per_paper.idxmax()
130     highest_avg_rating = avg_rating_per_paper.max()
131     print(f"Paper(s) with the highest average rating (ID: Average Rating):")
132     # Handle cases where multiple papers have the same highest rating
133     highestRatedPapers = avg_rating_per_paper[avg_rating_per_paper == highest_avg_rating]
134     print(highestRatedPapers)
135 else:
136     print("No papers with valid ratings found.")
137 print("-" * 40)
138
```

PROBLEMS   OUTPUT   DEBUG CONSOLE   TERMINAL   PORTS

Problem 6: Calculate the average rating given across all reviews (excluding NaNs).

Solution:

Average rating: 3.47

-----

Problem 7: Find the paper with the highest average rating.

Solution:

Paper(s) with the highest average rating (ID: Average Rating):

paper\_id

103     3.75

Name: rating, dtype: float64

## PROBLEM STATEMENTS AND ANSWERS 8 AND 9:

C: > Users > Atulya > OneDrive > Desktop > test.py > ...

```
139 # Problem 8: Determine the distribution of recommendations ('Accept', 'Reject', etc.).
140 print("\nProblem 8: Determine the distribution of recommendations.")
141 print("Solution:")
142 recommendation_distribution = df['recommendation'].value_counts()
143 print("Distribution of recommendations:")
144 print(recommendation_distribution)
145 print("-" * 40)
146
147 # Problem 9: Filter the dataset to show only reviews with a rating of 5.
148 print("\nProblem 9: Filter the dataset to show only reviews with a rating of 5.")
149 print("Solution:")
150 rating_5_reviews = df[df['rating'] == 5]
151 print("Reviews with a rating of 5:")
152 print(rating_5_reviews)
153 print("-" * 40)
```

PROBLEMS   OUTPUT   DEBUG CONSOLE   TERMINAL   PORTS

Problem 8: Determine the distribution of recommendations.

Solution:

Distribution of recommendations:

recommendation

Accept            9

Minor Revision   5

Major Revision   3

Reject            3

Name: count, dtype: int64

-----

Problem 9: Filter the dataset to show only reviews with a rating of 5.

Solution:

Reviews with a rating of 5:

|    | review_id | paper_id | reviewer_id | ... | recommendation | reviewer_expertise | review_length |
|----|-----------|----------|-------------|-----|----------------|--------------------|---------------|
| 7  | 8         | 104      | 204         | ... | Accept         | NaN                | 44            |
| 15 | 16        | 104      | 204         | ... | Accept         | High               | 30            |
| 18 | 19        | 103      | 201         | ... | Accept         | Medium             | 21            |

[3 rows x 9 columns]

-----

## PROBLEM STATEMENTS AND ANSWERS 10 AND 11:

```
C: > Users > Atulya > OneDrive > Desktop > test.py > ...
155 # Problem 10: Find the average rating for each reviewer expertise level.
156 print("\nProblem 10: Find the average rating for each reviewer expertise level.")
157 print("Solution:")
158 avg_rating_by_expertise = df.groupby('reviewer_expertise')['rating'].mean()
159 print("Average rating by reviewer expertise level:")
160 print(avg_rating_by_expertise)
161 print("-" * 40)
162
163 # Problem 11: Identify the reviewer who has submitted the most reviews.
164 print("\nProblem 11: Identify the reviewer who has submitted the most reviews.")
165 print("Solution:")
166 # Count reviews per reviewer and find the reviewer with the max count
167 reviews_per_reviewer = df['reviewer_id'].value_counts()
168 if not reviews_per_reviewer.empty:
169     most_active_reviewer_id = reviews_per_reviewer.idxmax()
170     most_reviews_count = reviews_per_reviewer.max()
171     print(f"Reviewer who submitted the most reviews: Reviewer ID {most_active_reviewer_id} ({most_reviews_count} reviews)")
172 else:
173     print("No reviews found.")
174 print("-" * 40)
175
```

PROBLEMS   OUTPUT   DEBUG CONSOLE   TERMINAL   PORTS

Problem 10: Find the average rating for each reviewer expertise level.  
Solution:  
Average rating by reviewer expertise level:  
reviewer\_expertise  
High      4.000000  
Low        3.000000  
Medium    3.142857  
Name: rating, dtype: float64  
-----

Problem 11: Identify the reviewer who has submitted the most reviews.  
Solution:  
Reviewer who submitted the most reviews: Reviewer ID 201 (5 reviews)  
-----

## PROBLEM STATEMENTS AND ANSWERS 12 AND 13:

```
C: > Users > Atulya > OneDrive > Desktop > test.py > ...
175
176 # Problem 12: Calculate the length of each review text (number of characters) and add it as a new column 'review_length'.
177 # Note: This was already done during dataset creation for convenience, but the code is shown here.
178 print("\nProblem 12: Calculate the length of each review text and add it as a new column 'review_length'.")
179 print("Solution: (column 'review_length' already added during setup)")
180 # df['review_length'] = df['review_text'].str.len() # This line would add the column if it didn't exist
181 print(df[['review_text', 'review_length']].head())
182 print("-" * 40)
183
184 # Problem 13: Find the average review length.
185 print("\nProblem 13: Find the average review length.")
186 print("Solution:")
187 average_review_length = df['review_length'].mean()
188 print(f"Average review length (character count): {average_review_length:.2f}")
189 print("-" * 40)
190
```

PROBLEMS   OUTPUT   DEBUG CONSOLE   TERMINAL   PORTS

-----

Problem 12: Calculate the length of each review text and add it as a new column 'review\_length'.  
Solution: (column 'review\_length' already added during setup)

|   | review_text                                | review_length |
|---|--|---------------|
| 0 | Good paper, clear methodology.             | 30            |
| 1 | Needs significant improvements in writing. | 42            |
| 2 | Excellent contribution to the field.       | 36            |
| 3 | Minor issues with figures, otherwise good. | 42            |
| 4 | Acceptable work, but lacks novelty.        | 35            |

-----

Problem 13: Find the average review length.  
Solution:  
Average review length (character count): 34.95  
-----

## PROBLEM STATEMENT AND ANSWER 14:

```
test.py x
C: > Users > Atulya > OneDrive > Desktop > test.py > ...

190
191 # Problem 14: Determine the paper(s) that received the most reviews.
192 print("\nProblem 14: Determine the paper(s) that received the most reviews.")
193 print("Solution:")
194 reviews_count_per_paper = df['paper_id'].value_counts()
195 if not reviews_count_per_paper.empty:
196     most_reviewed_paper_id = reviews_count_per_paper.idxmax()
197     most_reviews_count = reviews_count_per_paper.max()
198     print(f"Paper(s) that received the most reviews (ID: Number of Reviews):")
199     # Handle cases where multiple papers received the same highest number of reviews
200     most_reviewed_papers = reviews_count_per_paper[reviews_count_per_paper == most_reviews_count]
201     print(most_reviewed_papers)
202 else:
203     print("No papers found.")
204 print("-" * 40)
205
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
-----
Problem 14: Determine the paper(s) that received the most reviews.
Solution:
Paper(s) that received the most reviews (ID: Number of Reviews):
paper_id
101      5
Name: count, dtype: int64
-----
```



## PROBLEM STATEMENT AND ANSWER 15:

C: > Users > Atulya > OneDrive > Desktop > test.py > ...

```
206 # Problem 15: Filter reviews submitted after a specific date (e.g., '2023-03-01').
207 print("\nProblem 15: Filter reviews submitted after a specific date (e.g., '2023-03-01').")
208 print("Solution:")
209 # Ensure 'submission_date' is in datetime format (already done during setup)
210 # df['submission_date'] = pd.to_datetime(df['submission_date'])
211 reviews_after_date = df[df['submission_date'] > '2023-03-01']
212 print("Reviews submitted after 2023-03-01:")
213 print(reviews_after_date)
214 print("-" * 40)
215
```

PROBLEMS    OUTPUT    DEBUG CONSOLE    TERMINAL    PORTS

Problem 15: Filter reviews submitted after a specific date (e.g., '2023-03-01').

Solution:

Reviews submitted after 2023-03-01:

|    | review_id | paper_id | reviewer_id | ... | recommendation | reviewer_expertise | review_length |
|----|-----------|----------|-------------|-----|----------------|--------------------|---------------|
| 8  | 9         | 102      | 204         | ... | Accept         | Low                | 40            |
| 9  | 10        | 104      | 203         | ... | Reject         | Medium             | 22            |
| 10 | 11        | 105      | 205         | ... | Major Revision | High               | 43            |
| 11 | 12        | 101      | 201         | ... | Accept         | High               | 45            |
| 12 | 13        | 103      | 202         | ... | Minor Revision | Medium             | 40            |
| 13 | 14        | 105      | 205         | ... | Reject         | Low                | 27            |
| 14 | 15        | 102      | 203         | ... | Minor Revision | NaN                | 33            |
| 15 | 16        | 104      | 204         | ... | Accept         | High               | 30            |
| 16 | 17        | 101      | 202         | ... | Minor Revision | Medium             | 24            |
| 17 | 18        | 105      | 205         | ... | Major Revision | High               | 39            |
| 18 | 19        | 103      | 201         | ... | Accept         | Medium             | 21            |
| 19 | 20        | 104      | 203         | ... | Reject         | Low                | 34            |

[12 rows x 9 columns]

-----

## PROBLEM STATEMENTS AND ANSWERS 16 AND 17:

```
C: > Users > Atulya > OneDrive > Desktop > test.py > ...
216 # Problem 16: Count the number of reviews for each paper.
217 print("\nProblem 16: Count the number of reviews for each paper.")
218 print("Solution:")
219 reviews_count_per_paper_sorted = df['paper_id'].value_counts().sort_index()
220 print("Number of reviews for each paper:")
221 print(reviews_count_per_paper_sorted)
222 print("-" * 40)
223
224 # Problem 17: Find the minimum and maximum rating given (excluding NaNs).
225 print("\nProblem 17: Find the minimum and maximum rating given (excluding NaNs).")
226 print("Solution:")
227 min_rating = df['rating'].min()
228 max_rating = df['rating'].max()
229 print(f"Minimum rating: {min_rating}, Maximum rating: {max_rating}")
230 print("-" * 40)
231
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

-----  
Problem 16: Count the number of reviews for each paper.

Solution:

Number of reviews for each paper:

paper\_id

101 5

102 4

103 4

104 4

105 3

Name: count, dtype: int64  
-----

Problem 17: Find the minimum and maximum rating given (excluding NaNs).

Solution:

Minimum rating: 2.0, Maximum rating: 5.0  
-----

## PROBLEM STATEMENT AND ANSWER 18:

```
C: > Users > Atulya > OneDrive > Desktop > test.py > ...
232 # Problem 18: Group reviews by paper and find the average rating and the number of reviews for each paper.
233 print("\nProblem 18: Group reviews by paper and find the average rating and the number of reviews for each paper.")
234 print("Solution:")
235 paper_summary = df.groupby('paper_id').agg(
236     average_rating=('rating', 'mean'),
237     number_of_reviews=('review_id', 'count')
238 )
239 print("Summary (average rating and review count) for each paper:")
240 print(paper_summary)
241 print("-" * 40)
242
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

-----  
Problem 18: Group reviews by paper and find the average rating and the number of reviews for each paper.

Solution:

Summary (average rating and review count) for each paper:

average\_rating number\_of\_reviews

paper\_id

101

3.50

5

102

3.50

4

103

3.75

4

104

3.50

4

105

3.00

3  
-----

## PROBLEM STATEMENT AND ANSWER 19:

```
C: > Users > Atulya > OneDrive > Desktop > test.py > ...

243 # Problem 19: Check for missing values in each column.
244 print("\nProblem 19: Check for missing values in each column.")
245 print("Solution:")
246 missing_values = df.isnull().sum()
247 print("Missing values per column:")
248 print(missing_values)
249 print("-" * 40)

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

Problem 19: Check for missing values in each column.
Solution:
Missing values per column:
review_id      0
paper_id       0
reviewer_id    0
review_text    0
rating         3
submission_date 0
recommendation 0
reviewer_expertise 2
review_length  0
dtype: int64
```

## PROBLEM STATEMENT AND ANSWER 20:

```
C: > Users > Atulya > OneDrive > Desktop > test.py > ...

251 # Problem 20: Replace any missing values in the 'rating' column with the average rating.
252 print("\nProblem 20: Replace any missing values in the 'rating' column with the average rating.")
253 print("Solution:")
254 # Calculate the mean rating excluding NaNs for filling
255 mean_rating_for_fillna = df['rating'].mean()
256 df['rating'].fillna(mean_rating_for_fillna, inplace=True)
257 print(f"Missing values in 'rating' column after filling with mean ({mean_rating_for_fillna:.2f}):")
258 print(df['rating'].isnull().sum())
259 print("\nFirst 5 rows of 'rating' column after filling NaNs:")
260 print(df['rating'].head())
261 print("-" * 40)

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

Problem 20: Replace any missing values in the 'rating' column with the average rating.
Solution:
C:\Users\Atulya\OneDrive\Desktop\test.py:258: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
0

First 5 rows of 'rating' column after filling NaNs:
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
0

df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
0
```

C: > Users > Atulya > OneDrive > Desktop > test.py > ...

251 # Problem 20: Replace any missing values in the 'rating' column with the average rating

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
0
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
0
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
0
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
df['rating'].fillna(mean_rating_for_fillna, inplace=True)
Missing values in 'rating' column after filling with mean (3.47):
0
Missing values in 'rating' column after filling with mean (3.47):
Missing values in 'rating' column after filling with mean (3.47):
0
```



```
First 5 rows of 'rating' column after filling NaNs:
```

```
0    4.000000
```

```
First 5 rows of 'rating' column after filling NaNs:
```

```
0    4.000000
```

```
1    2.000000
```

```
2    3.470588
```

```
0    4.000000
```

```
1    2.000000
```

```
2    3.470588
```

```
1    2.000000
```

```
2    3.470588
```

```
2    3.470588
```

```
3    3.000000
```

```
4    4.000000
```

```
Name: rating, dtype: float64
```

```
3    3.000000
```

```
4    4.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
4    4.000000
```

```
Name: rating, dtype: float64
```

```
-----
```

```
3    3.000000
```

```
4    4.000000
```

```
Name: rating, dtype: float64
```

```
3    3.000000
```

```
4    4.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
4    4.000000
```

```
Name: rating, dtype: float64
```

```
-----  
PS C:\Users\Atulya\OneDrive\Desktop>
```

```
3    3.000000
```

```
4    4.000000
```

```
Name: rating, dtype: float64
```

```
-----  
PS C:\Users\Atulya\OneDrive\Desktop>
```

```
3    3.000000
```

```
4    4.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
3    3.000000
```

```
4    4.000000
```

```
Name: rating, dtype: float64
```

```
-----  
PS C:\Users\Atulya\OneDrive\Desktop> █
```